



# Distributed deep learning and why you may not need it

Jakub Sanojca, Mikuláš Zelinka

PyData Warsaw 2018

# About us


- **Research at Avast**
  - **400M customers (Avast, AVG, HideMyAss, CCleaner)**
  - **Tons of ~~damage~~ data**



# Outline

- **Wow such Deep Learning much AI**
- **Different kinds of distributed learning and parallelism**
- **Frameworks and approaches**

# One picture explaining the rise of Deep Learning




Andrew Ng

# Revisiting Unreasonable Effectiveness of Data in Deep Learning Era


Chen Sun, Abhinav Shrivastava, Saurabh Singh, and Abhinav Gupta

- COCO object detection:




# Why go distributed?

Demystifying Parallel and Distributed Deep Learning: An In-Depth Concurrency Analysis 1:7




(a) Hardware Architectures




(b) Training with Single vs. Multiple Nodes

# Types of parallelism

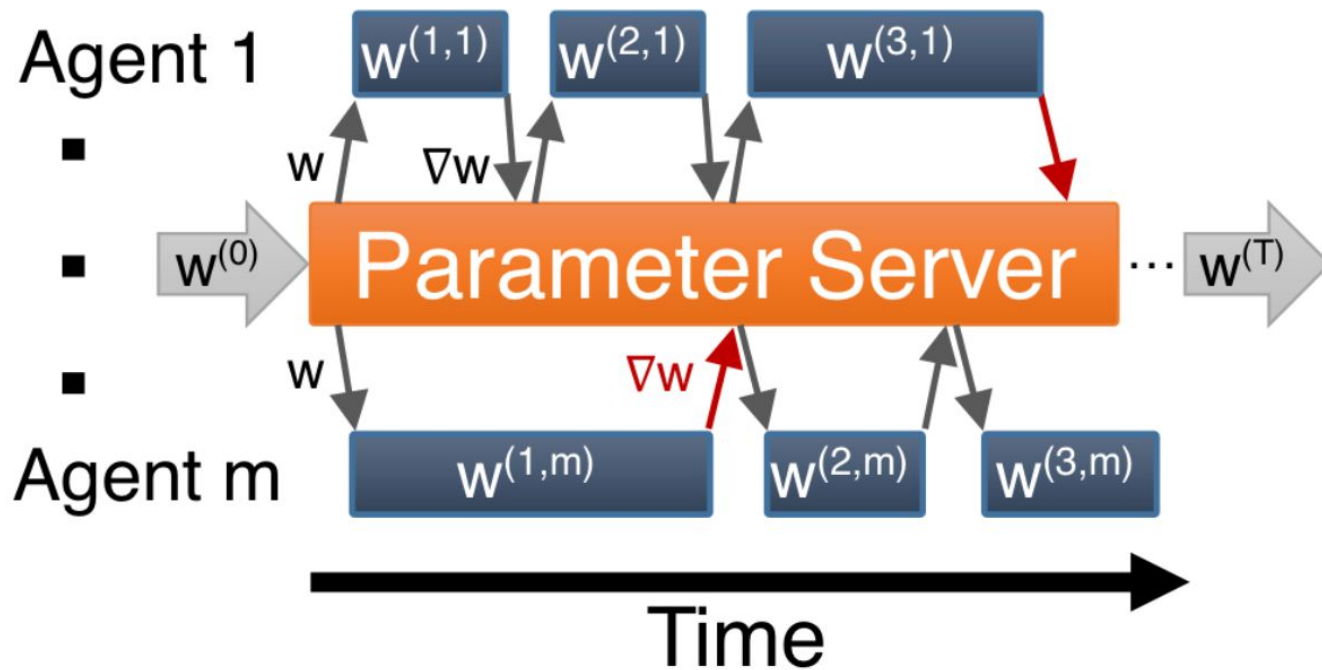
Model Parallelism



Data Parallelism

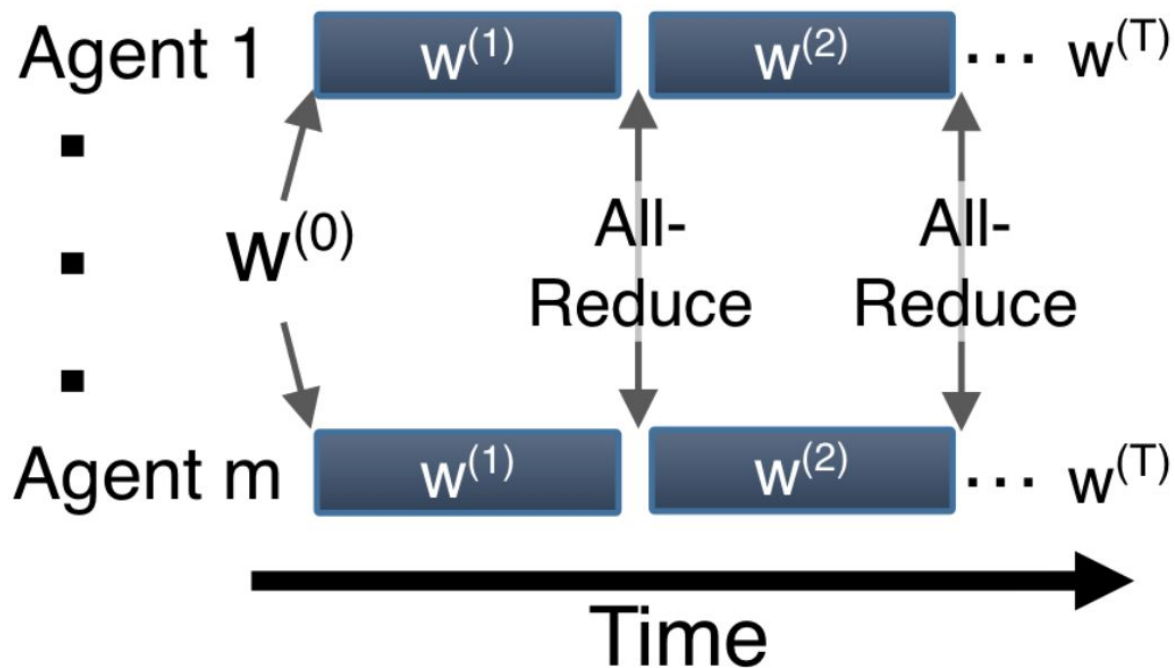


# Asynchronous training







# Synchronous training with allreduce



# Model parallelism





# Distributed Tensorflow (asynchronous)

```
# Create a cluster from the parameter server and worker hosts.
cluster = tf.train.ClusterSpec({"ps": ps_hosts, "worker": worker_hosts})

# Create and start a server for the local task.
server = tf.train.Server(cluster,
                        job_name=FLAGS.job_name,
                        task_index=FLAGS.task_index)

if FLAGS.job_name == "ps":
    server.join()
elif FLAGS.job_name == "worker":

    # Assigns ops to the local worker by default.
    with tf.device(tf.train.replica_device_setter(
        worker_device="/job:worker/task:%d" % FLAGS.task_index,
        cluster=cluster)):

        # Build model...

# The MonitoredTrainingSession takes care of session initialization,
# restoring from a checkpoint, saving to a checkpoint, and closing when done
# or an error occurs.
    with tf.train.MonitoredTrainingSession(master=server.target,
        is_chief=(FLAGS.task_index == 0),
        checkpoint_dir="/tmp/train_logs",
        hooks=hooks) as mon_sess:

        # Train your model
```



# Tensorflow

```
run_config = tf.estimator.RunConfig()

classifier = tf.estimator.Estimator(
    model_fn=model_function, # description of your model - loss, updates etc.
    model_dir=model_dir, # directory where you want to store your model
    config=run_config) # how often you need summaries, how often to checkpoint etc.

classifier.train(input_fn=input_function)
```

# Distributed Tensorflow (synchronous, single node)

```
distribution = tf.contrib.distribute.MirroredStrategy() # multi-GPU distribution

run_config = tf.estimator.RunConfig(train_distribute=distribution)

classifier = tf.estimator.Estimator(
    model_fn=model_function, # description of your model - loss, updates etc.
    model_dir=model_dir, # directory where you want to store your model
    config=run_config) # how often you need summaries, how often to checkpoint etc.

classifier.train(input_fn=input_function)
```



# Horovod



```
import tensorflow as tf
import horovod.tensorflow as hvd

# Initialize Horovod
hvd.init()

# Pin GPU to be used to process local rank (one GPU per process)
config = tf.ConfigProto()
config.gpu_options.visible_device_list = str(hvd.local_rank())

# Build model...

# Add Horovod Distributed Optimizer
opt = hvd.DistributedOptimizer(opt)

# The MonitoredTrainingSession takes care of session initialization,
# restoring from a checkpoint, saving to a checkpoint, and closing when done
# or an error occurs.
with tf.train.MonitoredTrainingSession(checkpoint_dir=checkpoint_dir,
                                       config=config,
                                       hooks=hooks) as mon_sess:

    # Perform synchronous training.
```

# Horovod



```
import tensorflow as tf
import horovod.tensorflow as hvd

# Initialize Horovod
hvd.init()

# Pin GPU to be used to process local rank (one GPU per process)
config = tf.ConfigProto()
config.gpu_options.visible_device_list = str(hvd.local_rank())

# Build model...

# Add Horovod Distributed Optimizer
opt = hvd.DistributedOptimizer(opt)

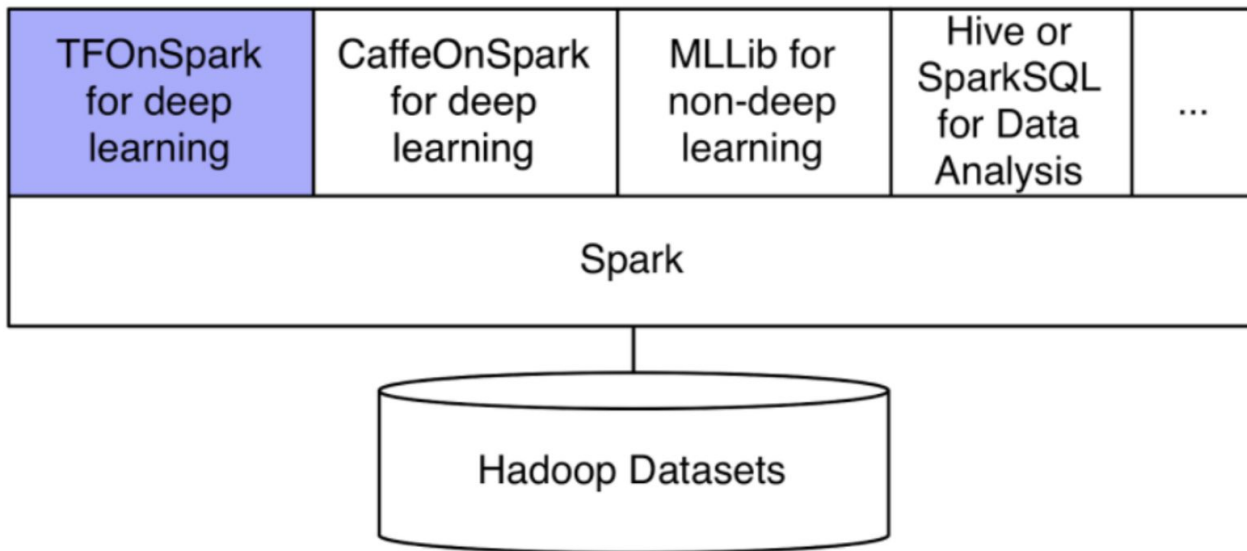
# The MonitoredTrainingSession takes care of session initialization,
# restoring from a checkpoint, saving to a checkpoint, and closing when done
# or an error occurs.
with tf.train.MonitoredTrainingSession(checkpoint_dir=checkpoint_dir,
                                       config=config,
                                       hooks=hooks) as mon_sess:

    # Perform synchronous training.
```

```
$ mpirun -np 16 \
  -H server1:4,server2:4,server3:4,server4:4 \
  -bind-to none -map-by slot \
  -x NCCL_DEBUG=INFO -x LD_LIBRARY_PATH -x PATH \
  -mca pml ob1 -mca btl ^openib \
  python train.py
```




# Tensorflow on Spark




# Tensorflow on Spark

## InputMode.SPARK



# Tensorflow on Spark

## InputMode.TENSORFLOW



# Tensorflow on Spark

```
from tensorflowonspark import TFCluster, TFNode

# start the cluster
cluster = TFCluster.run(sc, map_fn, args, num_executors, num_ps, tensorboard, input_mode)

# either pass data for training
cluster.train(dataRDD, num_epochs=0)

# or pass data for scoring
cluster.inference(dataRDD)

# shutdown the cluster
cluster.shutdown()
```

# Tensorflow on Spark

```
def map_fun(args, ctx):  
  
    import tensorflow as tf  
  
    # necessary variables  
    logdir = ctx.absolute_path(args.model)  
    hooks = [tf.train.StopAtStepHook(last_step=100000)]  
    batch_size = args.batch_size  
  
    worker_num = ctx.worker_num  
    job_name = ctx.job_name  
    task_index = ctx.task_index  
  
    # Get TF cluster and server instances  
    cluster, server = ctx.start_cluster_server(1, args.rdma)  
  
    if job_name == "ps":  
        print('PARAMETER SERVER')  
        server.join()  
    elif job_name == "worker":  
  
        # Assigns ops to workers - taken care of by Spark  
        with tf.device(tf.train.replica_device_setter(  
            worker_device="/job:worker/task:%d" % task_index,  
            cluster=cluster)):  
  
            [...]   
  
        # start your MonitoredTrainingSession  
        with tf.train.MonitoredTrainingSession(master=server.target,  
                                                is_chief=(task_index == 0),  
                                                checkpoint_dir=logdir,  
                                                hooks=hooks) as mon_sess:  
  
            # train or score your data  
            [...]
```

# Comparison and our benchmarks

<b>synchronous Tensorflow</b>	<b>asynchronous Tensorflow</b>	<b>Horovod</b>	<b>TFonS</b>
easiest to use	not hard to use, but requires code change	requires code change	requires most code change
only one machine	multi-device, multi-machine	multi-device, multi-machine	multi-device, multi-machine
speed scales almost linearly	scales almost linearity	slower than expected, most probably because of higher network requirements	slowest because of a lot of overhead, can be run on existing infrastructure

# Why you may not need it

- **Hardware improvements**
- **Code simplicity**
- **Simpler models are enough / data efficiency**
- **Experimentation**



**Thank you!**

**For questions, feel free to reach out to us at:**

[jakub.sanojca@avast.com](mailto:jakub.sanojca@avast.com)

[mikulas.zelinka@avast.com](mailto:mikulas.zelinka@avast.com)



# Resources:

- <https://arxiv.org/pdf/1802.09941.pdf> / <https://www.youtube.com/watch?v=xtxxLWZznBI>
- <https://arxiv.org/pdf/1610.02527.pdf>
- <https://papers.nips.cc/paper/4687-large-scale-distributed-deep-networks.pdf>
- <https://medium.com/@esaliya/model-parallelism-in-deep-learning-is-not-what-you-think-94d2f81e82ed>
- <https://blog.skymind.ai/distributed-deep-learning-part-1-an-introduction-to-distributed-training-of-neural-networks/>
- <https://www.youtube.com/watch?v=SphfeTI70MI>
- <https://www.tensorflow.org/deploy/distributed>
- <https://arxiv.org/pdf/1802.05799.pdf> / <https://github.com/uber/horovod>
- <https://github.com/yahoo/TensorFlowOnSpark>
- <https://github.com/tensorflow/ecosystem/tree/master/spark/spark-tensorflow-connector>